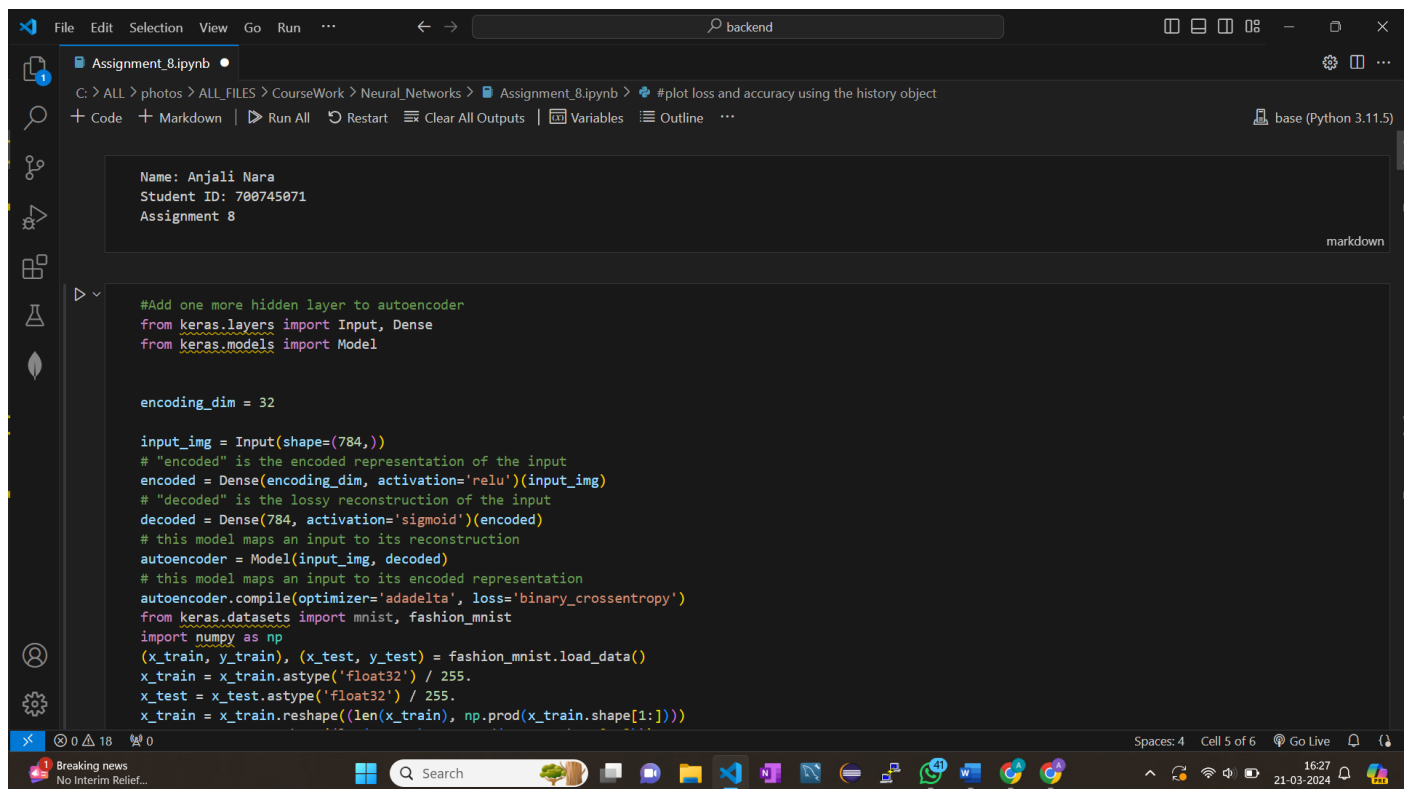


Video Link: https://drive.google.com/file/d/1R_daHBknReThxbpqi78IfFrDac5k4RbD/view?usp=sharing

Github Link: <https://github.com/naraanjali/Neural-assignment-8/tree/main>

1. Add one more hidden layer to autoencoder



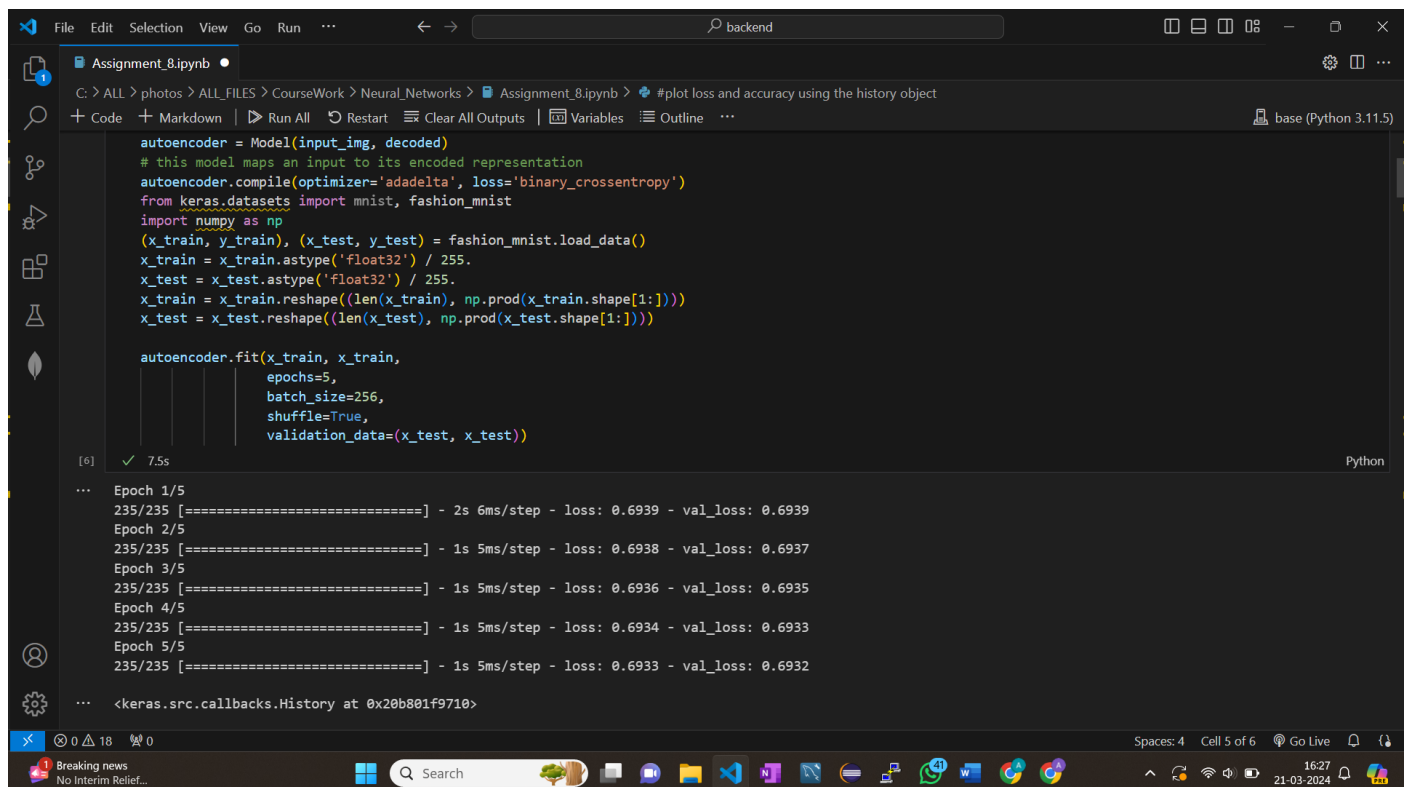
The screenshot shows a Jupyter Notebook titled 'Assignment_8.ipynb'. The code defines an autoencoder with one hidden layer. The input image is flattened to a vector of size 784. The encoder consists of a Dense layer with 32 units and 'relu' activation. The decoder consists of a Dense layer with 784 units and 'sigmoid' activation. The model is compiled with the 'adadelta' optimizer and 'binary_crossentropy' loss. The training data is loaded from the Fashion-MNIST dataset and reshaped to (len(x_train), np.prod(x_train.shape[1:])).

```
Name: Anjali Nara
Student ID: 700745071
Assignment 8

#Add one more hidden layer to autoencoder
from keras.layers import Input, Dense
from keras.models import Model

encoding_dim = 32

input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
# this model maps an input to its reconstruction
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
from keras.datasets import mnist, fashion_mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
```



The screenshot shows the continuation of the Jupyter Notebook. The code defines the autoencoder model and compiles it. The training process is shown, with the model being trained for 5 epochs. The training and validation loss are displayed for each epoch. The model is trained on the Fashion-MNIST dataset.

```
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
from keras.datasets import mnist, fashion_mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

autoencoder.fit(x_train, x_train,
                epochs=5,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test))
```

[6] ✓ 7.5s

```
... Epoch 1/5
235/235 [=====] - 2s 6ms/step - loss: 0.6939 - val_loss: 0.6939
Epoch 2/5
235/235 [=====] - 1s 5ms/step - loss: 0.6938 - val_loss: 0.6937
Epoch 3/5
235/235 [=====] - 1s 5ms/step - loss: 0.6936 - val_loss: 0.6935
Epoch 4/5
235/235 [=====] - 1s 5ms/step - loss: 0.6934 - val_loss: 0.6933
Epoch 5/5
235/235 [=====] - 1s 5ms/step - loss: 0.6933 - val_loss: 0.6932
... <keras.src.callbacks.History at 0x20b801f9710>
```

2. Do the prediction on the test data and then visualize one of the reconstructed version of that test data. Also, visualize the same test data before reconstruction using Matplotlib

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline ... base (Python 3.11.5)

# Do the prediction on the test data and then visualize one of the reconstructed version of that test data.
#Also, visualize the same test data before reconstruction using Matplotlib

from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist, fashion_mnist
import numpy as np
import matplotlib.pyplot as plt

encoding_dim = 32

input_img = Input(shape=(784,))

hidden_1 = Dense(256, activation='relu')(input_img)

encoded = Dense(encoding_dim, activation='relu')(hidden_1)

hidden_2 = Dense(256, activation='relu')(encoded)

# Define the output layer
decoded = Dense(784, activation='sigmoid')(hidden_2)

# Define the autoencoder model
autoencoder = Model(input_img, decoded)

# Compile the model
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()
```

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline ... base (Python 3.11.5)

# Compile the model
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()

x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

history = autoencoder.fit(x_train, x_train,
                        epochs=5,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test, x_test))

decoded_imgs = autoencoder.predict(x_test)

# Visualize one of the reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original test image
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
```

Assignment_8.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object

+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline

base (Python 3.11.5)

```
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

# Display reconstructed test image
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(decoded_imgs[i].reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()
```

[7] ✓ 14.2s Python

Epoch 1/5

0 18 0

Breaking news
No Interim Relief...

Search

Spaces: 4 Cell 5 of 6 Go Live 16:29 21-03-2024

Assignment_8.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object

+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline

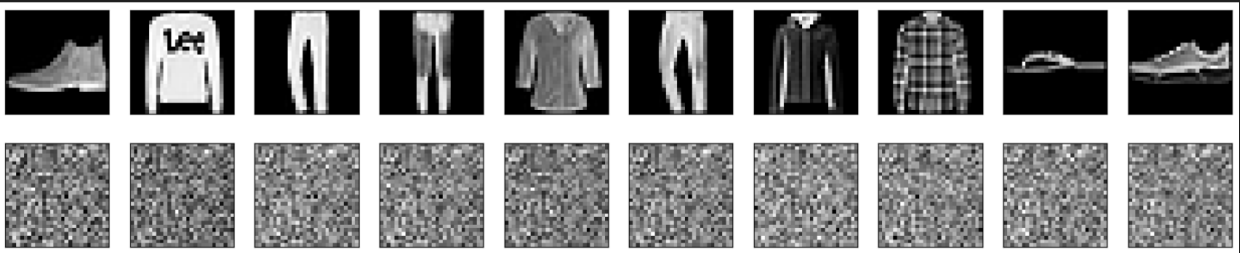
base (Python 3.11.5)

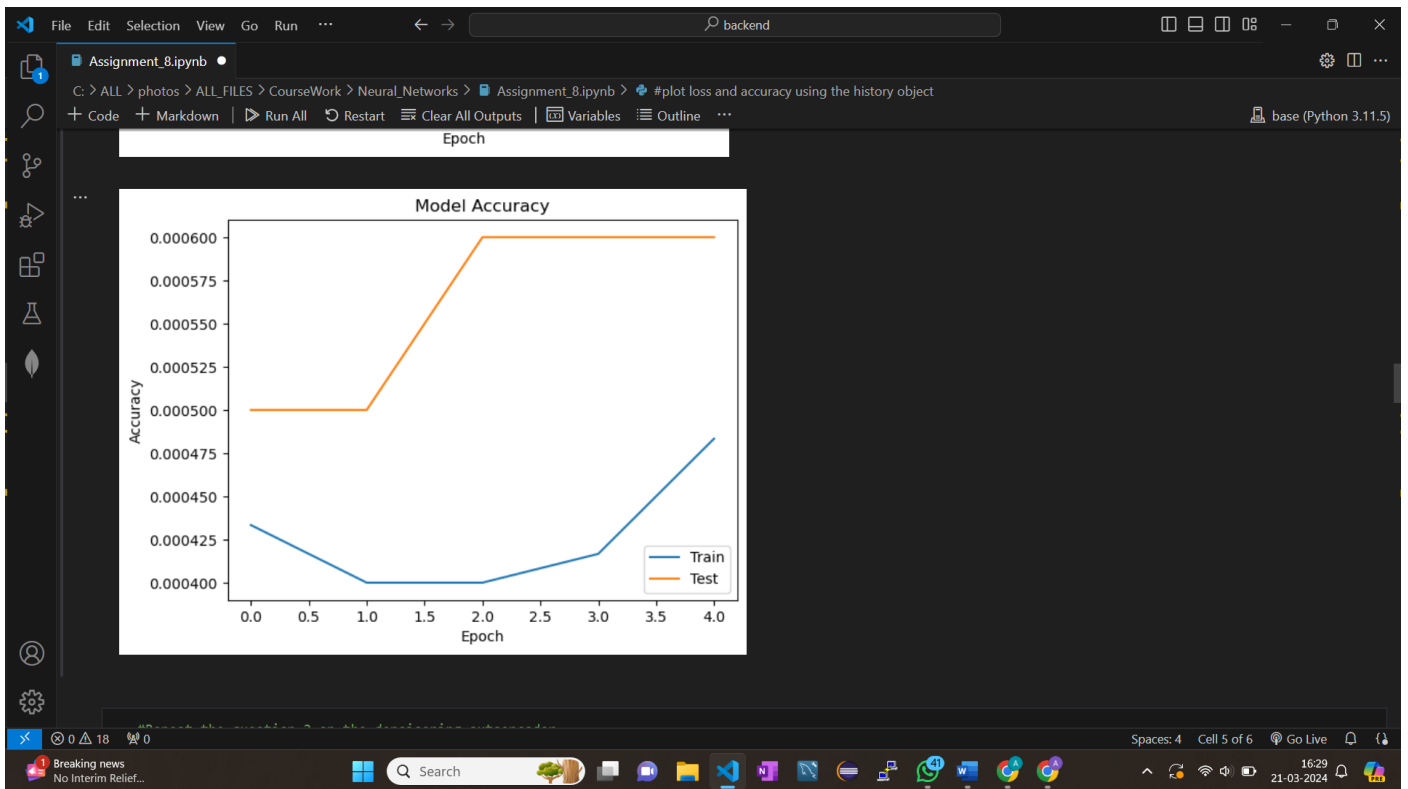
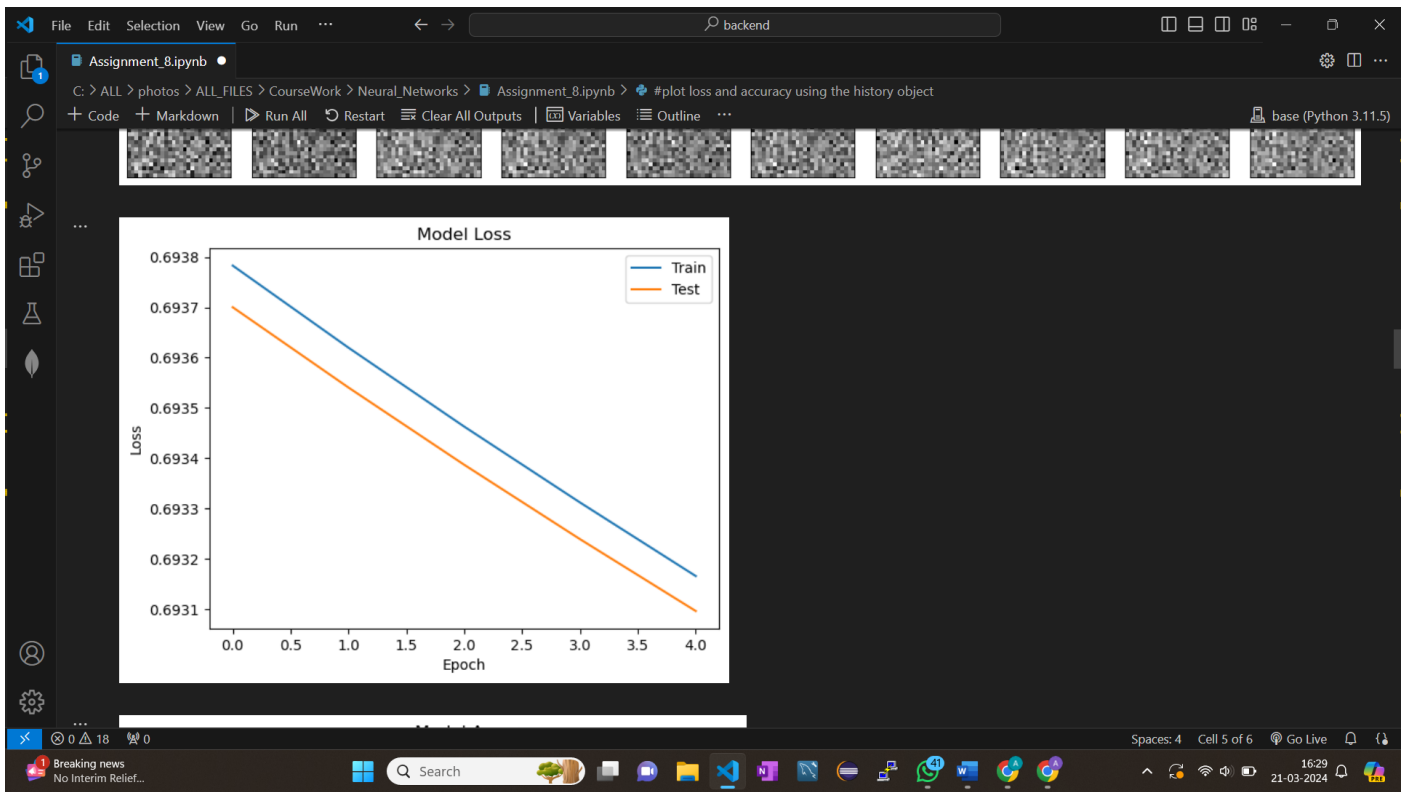
```
plt.show()
```

[7] ✓ 14.2s Python

Epoch 1/5
235/235 [=====] - 3s 12ms/step - loss: 0.6938 - accuracy: 4.3333e-04 - val_loss: 0.6937 - val_accuracy: 5.0000e-04
Epoch 2/5
235/235 [=====] - 2s 10ms/step - loss: 0.6936 - accuracy: 4.0000e-04 - val_loss: 0.6935 - val_accuracy: 5.0000e-04
Epoch 3/5
235/235 [=====] - 2s 9ms/step - loss: 0.6935 - accuracy: 4.0000e-04 - val_loss: 0.6934 - val_accuracy: 6.0000e-04
Epoch 4/5
235/235 [=====] - 2s 9ms/step - loss: 0.6933 - accuracy: 4.1667e-04 - val_loss: 0.6932 - val_accuracy: 6.0000e-04
Epoch 5/5
235/235 [=====] - 2s 9ms/step - loss: 0.6932 - accuracy: 4.8333e-04 - val_loss: 0.6931 - val_accuracy: 6.0000e-04
313/313 [=====] - 1s 1ms/step

...





3. Repeat the question 2 on the denoising autoencoder

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline ...
base (Python 3.11.5)

#Repeat the question 2 on the denoising autoencoder
from keras.layers import Input, Dense
from keras.models import Model

encoding_dim = 32 # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

input_img = Input(shape=(784,))
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input_img, decoded)
# this model maps an input to its encoded representation
autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
from keras.datasets import fashion_mnist
import numpy as np
(x_train, _), (x_test, _) = fashion_mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

noise_factor = 0.5
x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

autoencoder.fit(x_train_noisy, x_train,
                epochs=10,
                batch_size=256,
                shuffle=True,
```

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown | Run All | Restart | Clear All Outputs | Variables | Outline ...
base (Python 3.11.5)

autoencoder.fit(x_train_noisy, x_train,
                epochs=10,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test_noisy, x_test_noisy))

[0] ✓ 14.5s Python

... Epoch 1/10
235/235 [=====] - 2s 6ms/step - loss: 0.6977 - val_loss: 0.6977
Epoch 2/10
235/235 [=====] - 1s 5ms/step - loss: 0.6973 - val_loss: 0.6973
Epoch 3/10
235/235 [=====] - 1s 5ms/step - loss: 0.6970 - val_loss: 0.6970
Epoch 4/10
235/235 [=====] - 1s 5ms/step - loss: 0.6967 - val_loss: 0.6967
Epoch 5/10
235/235 [=====] - 1s 5ms/step - loss: 0.6965 - val_loss: 0.6965
Epoch 6/10
235/235 [=====] - 1s 5ms/step - loss: 0.6962 - val_loss: 0.6962
Epoch 7/10
235/235 [=====] - 1s 5ms/step - loss: 0.6959 - val_loss: 0.6959
Epoch 8/10
235/235 [=====] - 1s 6ms/step - loss: 0.6956 - val_loss: 0.6957
Epoch 9/10
235/235 [=====] - 1s 6ms/step - loss: 0.6954 - val_loss: 0.6954
Epoch 10/10
235/235 [=====] - 1s 6ms/step - loss: 0.6951 - val_loss: 0.6952

... <keras.src.callbacks.History at 0x20bc8b19ed0>
```

4. plot loss and accuracy using the history object

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown + Run All + Restart + Clear All Outputs + Variables + Outline ...
base (Python 3.11.5)

#plot loss and accuracy using the history object
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import fashion_mnist
import numpy as np
import matplotlib.pyplot as plt

encoding_dim = 32

input_img = Input(shape=(784,))

encoded = Dense(encoding_dim, activation='relu')(input_img)

decoded = Dense(784, activation='sigmoid')(encoded)

autoencoder = Model(input_img, decoded)

# Compile the model
autoencoder.compile(optimizer='adadelata', loss='binary_crossentropy', metrics=['accuracy'])

# Load the fashion MNIST dataset
(x_train, _), (x_test, _) = fashion_mnist.load_data()

# Normalize the data and flatten the images
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
```

```
Assignment_8.ipynb
C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object
+ Code + Markdown + Run All + Restart + Clear All Outputs + Variables + Outline ...
base (Python 3.11.5)

noise_factor = 0.5
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)

history = autoencoder.fit(x_train_noisy, x_train,
                        epochs=10,
                        batch_size=256,
                        shuffle=True,
                        validation_data=(x_test_noisy, x_test_noisy))

decoded_imgs = autoencoder.predict(x_test_noisy)

# Visualize one of the noisy test images
plt.figure(figsize=(20, 4))
n = 10
for i in range(n):
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test_noisy[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

# Visualize one of the reconstructed test images
for i in range(n):
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```

Assignment_8.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object

ax.get_yaxis().set_visible(False)

plt.show()

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper right')
plt.show()

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='lower right')
plt.show()

[9] ✓ 16.2s Python

Epoch 1/10
235/235 [=====] - 2s 7ms/step - loss: 0.6972 - accuracy: 0.0016 - val_loss: 0.6969 - val_accuracy: 0.0018
Epoch 2/10
235/235 [=====] - 1s 6ms/step - loss: 0.6967 - accuracy: 0.0016 - val_loss: 0.6964 - val_accuracy: 0.0019
Epoch 3/10
235/235 [=====] - 1s 6ms/step - loss: 0.6963 - accuracy: 0.0016 - val_loss: 0.6960 - val_accuracy: 0.0019
Epoch 4/10
235/235 [=====] - 1s 6ms/step - loss: 0.6959 - accuracy: 0.0016 - val_loss: 0.6956 - val_accuracy: 0.0020
Epoch 5/10
235/235 [=====] - 1s 6ms/step - loss: 0.6955 - accuracy: 0.0016 - val_loss: 0.6952 - val_accuracy: 0.0021

Assignment_8.ipynb

C: > ALL > photos > ALL_FILES > CourseWork > Neural_Networks > Assignment_8.ipynb > #plot loss and accuracy using the history object

235/235 [=====] - 1s 6ms/step - loss: 0.6963 - accuracy: 0.0016 - val_loss: 0.6960 - val_accuracy: 0.0019
Epoch 4/10
235/235 [=====] - 1s 6ms/step - loss: 0.6959 - accuracy: 0.0016 - val_loss: 0.6956 - val_accuracy: 0.0020
Epoch 5/10
235/235 [=====] - 1s 6ms/step - loss: 0.6955 - accuracy: 0.0016 - val_loss: 0.6952 - val_accuracy: 0.0021
Epoch 6/10
235/235 [=====] - 1s 6ms/step - loss: 0.6952 - accuracy: 0.0016 - val_loss: 0.6949 - val_accuracy: 0.0021
Epoch 7/10
235/235 [=====] - 1s 6ms/step - loss: 0.6948 - accuracy: 0.0016 - val_loss: 0.6945 - val_accuracy: 0.0019
Epoch 8/10
235/235 [=====] - 1s 6ms/step - loss: 0.6944 - accuracy: 0.0016 - val_loss: 0.6942 - val_accuracy: 0.0019
Epoch 9/10
235/235 [=====] - 1s 6ms/step - loss: 0.6941 - accuracy: 0.0016 - val_loss: 0.6938 - val_accuracy: 0.0020
Epoch 10/10
235/235 [=====] - 1s 6ms/step - loss: 0.6937 - accuracy: 0.0016 - val_loss: 0.6935 - val_accuracy: 0.0020
313/313 [=====] - 0s 1ms/step

...

